

NOMA v1

An agent-first connected-home app for Indian homes, anchored on an air purifier. Compiled from every source supplied to date.

Status: DRAFT — not approvable yet.

One scope decision (§3) and four critical blockers (§4) sit upstream of everything in this document. Thirteen conflicts are open; none has been resolved here, because the sources genuinely disagree and silently picking a winner would hide that.

VERSION

0.2 · 2026-08-04

DECISION LOG

memory/decisions.md · ADR-001...005

FEATURES REGISTERED

13 · per-feature PRDs: 0

CANONICAL VALUES

src/tokens · src/hardware

OPEN CONFLICTS

13 (4 critical)

ANALYTICS EVENTS

0

§ Contents

1. Sources and provenance	14. Data model
2. Problem and insight	15. Permissions and sharing
3. The scope decision	16. Privacy and compliance
4. Critical blockers	17. Competitive position
5. Users	18. Monetization
6. Jobs to be done	19. Voice and content
7. Product principles	20. Design system
8. The agent model	21. Motion
9. The Week One arc	22. Analytics and success metrics
10. Journeys	23. Risks
11. Feature scope	24. Open conflicts
12. The differentiator	25. Sequencing
13. Hardware and sensing	26. Declares

1 Sources and provenance

Everything in this document traces to something supplied. Where sources disagree, the disagreement is stated. Gaps are marked **GAP** rather than filled.

SOURCE	WHAT IT IS	FED INTO
Air Agent — agent-first sample UI (C2) 216KB · 2 versions	The most valuable document. The graduated-autonomy agent model: three home zones, per-capability autonomy tiers, the four-slot explanation contract, inert editable rehearsal, standing rules. The only document describing an <i>architecture</i> rather than screens.	Tech · User flows · Content
Ding Air Pro · Air Week One 60KB	A day-by-day first-week retention arc. Each day states its emotional target, the research finding behind it, and the screen that delivers it. The only document that argues rather than asserts.	Product · Journeys · Content · Metrics
Ding — Smart Home App (V2) 13MB · ~60 screens	The broadest build: onboarding, purifier, vacuum, camera, sharing, shop, score, wearable bridge. Zero static text — all content lives inside 739KB of JSX, extracted from ~45 data constants. Treat its naming as unreliable; it is the source of most conflicts.	Hardware · Data · Product
Noise Sanctuary — India's most trustworthy purifier 360KB	Positioning and commercial strategy: the AQI-weighted filter argument, a five-way competitor teardown, three subscription tiers, the whole-home ecosystem pitch.	Product · Monetization · Hardware
Airbnb iOS observed at 393×852pt	Visual reference for surface, radius and elevation only. Its warmth was taken; its density and its red were not.	Design (ADR-001)
Google Sans Flex · two green revisions · a motion ruling · a product render	Owner-supplied directly, not derived.	Design (ADR-002... 005)

PROVENANCE RISK

The four prototype files still live in the owner's `~/Downloads` . The provenance of this entire knowledge base is currently a local folder. They should be vendored into `docs/prototypes/` .

WHAT THESE SOURCES ARE NOT

Design prototypes, not specifications. Seeded data, happy paths only, and **no empty, loading, error or offline state anywhere across all four files.** That absence is systematic, not incidental — it is why the PRD template now mandates those states. Device specs (CADR, coverage, price) are marked `prototype-only` and have not been checked against a datasheet or BOM.

2 Problem and insight

You can't see what you breathe.

Air Week One

An air purifier is a box that hums. When it works, nothing visibly happens. The source material is blunt about the consequence:

Invisible protection that runs silently is, to the user, indistinguishable from a switched-off box.

Air Week One — the "away-mode trap"

In Indian homes this compounds. People switch purifiers *off* to save electricity — losing the benefit they paid for — because the cost is visible and the protection is not. 82% of Indians cannot see indoor air without a device.

The insight the product is built on

Air is invisible, so perceived value is whatever the app makes visible.

Perceived value is therefore a *design output*, not a hardware output. That single sentence justifies the entire seven-day arc in §9, the narration requirement in §8, and the existence of the energy card in §11. It is the most load-bearing line in the source material.

3 The scope decision

The supplied material describes two different products. Nothing says which ships first. This is the decision everything else waits on.

	A · AIR COMPANION	B · WHOLE-HOME PLATFORM
Anchor	One air purifier	Purifier + cameras + locks + vacuums + doorbells (+ lights, music on roadmap)
Evidence quality	Deep. A day-by-day retention arc with stated research findings, named behavioural mechanisms, and an identified chokepoint. An articulated agent architecture.	Broad but asserted. ~60 screens and a subscription deck, with no research behind the claims.
Revenue	Hardware + filters	Hardware + filters + ₹699/₹1,199 subscription
Open conflicts	6	13
Personas served	P1, P2, P4	All four — P3 (household help) only exists if there are locks and cameras

RECOMMENDATION

Ship A, architected for B. The purifier story is the only part of the material that argues rather than asserts. The platform story has more screens but less conviction, and carries most of the unresolved conflicts (roles, score, subscription, pricing). Build A's data model as a strict subset of B's so nothing is thrown away when the

ecosystem lands.

This is the owner's call. Everything below is scoped to A unless marked **B** . Until it is ruled, teams will build against different assumptions.

4 Critical blockers

Thirteen conflicts are open. These four are different in kind: each one stops work rather than slowing it. Full list in §24.

C-7 · THE DIFFERENTIATOR HAS NO ENGINEERING MODEL

AQI-weighted filter life is the single claim the product is built on, and it exists only as a worked example. The material states the effect ("above AQI 200 a filter absorbs 3–5× more particulate per hour") and shows Delhi at 38% vs Mumbai at 64% on identical runtime — but supplies **no formula, no coefficients, and no validation data**.

It gates `F-FILTER-HEALTH`, which is simultaneously the differentiator *and* the recurring-revenue lever. **Nothing else in v1 is worth as much as closing this.**

C-11 · "DATA NEVER LEAVES HOME" VS THREE NETWORKED FEATURES

The product claims on-device AI and DPDP compliance, then ships three features that cannot work locally: CPCB outdoor blending, forecast-driven pre-clean, and neighbourhood comparison. **"Never leaves home" and "uses tomorrow's forecast" cannot both be literally true.**

The claim most likely means *personal and biometric data stays local; ambient public data is fetched* — which is defensible and still a strong position. It must be narrowed to that before it reaches store copy or a privacy policy. It also blocks the analytics taxonomy, because an analytics SDK is by definition exfiltration.

C-8 · NO SAFETY CLASS FOR AGENT AUTONOMY

"Act silently" is offered per capability, but every example is comfort-grade (fan speed, pre-clean, reorder). Nothing prevents it applying to a **lock**, and `Reorder filters` at that level is an **unattended purchase**. Separately, the SpO_2 -below-94% → purifier-Turbo trigger is a health-adjacent inference presented as an ordinary automation, with no medical or legal position.

v1 ships "Ask me" only until capabilities carry a safety class.

C-2 · THE PRODUCT AND ASSISTANT HAVE FIVE NAMES BETWEEN THEM

Product: **NOMA** · Ding · Noise Home · Noise Air · Noise Sanctuary.

Assistant: Ding · Noise AI Agent · Sanctuary Assistant · unnamed first person.

Score: Home Score · Sanctuary Score · Air Score.

All user-facing copy, every i18n key, all persona prompts, the glossary and every SKU string are downstream. **Highest-leverage single decision on the list.**

5 Users

C-1 · PERSONAS UNRESOLVED

Two incompatible sets were supplied. `src/personas/personas.json` is deliberately empty. Set B is used throughout this PRD because the prototypes corroborate it four ways — **Ravi and Lakshmi appear as real named people in the camera event feed**, all four product surfaces Set B names are built (Home Score, Faces You Trust, Shop, AQI), P3's household help maps exactly onto the `staff` role, and **Set A's names appear nowhere in any of the four documents**. Strong evidence — still the owner's ruling.

Set B — role-based, used here

ID	PERSONA	NEEDS	DESIGN CONSEQUENCE
P1	Asha, 32 PM in Bangalore PRIMARY · BUYER · POWER USER	Manages her parents' separate home remotely. AQI-anxious, security-conscious, time-poor. Knows her family is safe without checking six apps; acts instantly when something is wrong.	She is the buyer <i>and</i> the power user, so she absorbs depth — but she is time-poor, so depth must be opt-in, never the default view .
P2	Ravi & Lakshmi , early 60s Living in the home Asha manages SECONDARY · RESIDENT	Lower app fluency. Want things to just work — door locks itself, air stays clean — with minimal interaction and large, legible controls.	Sets a hard floor on type size and touch targets . Any flow that <i>requires</i> app fluency to stay safe is a defect, not an edge case. A calm, sparse layout serves them directly.
P3	The Household Help Maid, cook, delivery agent TERTIARY · TIME-BOXED	Scoped, temporary entry. Should be recognised or granted limited access — never full account control.	Access is scoped and expiring by default . A permissions-model requirement, not a UI preference. Only exists under scope <code>B</code> .
P4	Priya Owns one device, considering more ACQUISITION	To understand what each device adds; to buy confidently inside an app she already trusts.	Monetisation is contextual and data-triggered. It competes with the primary action for the 10% accent budget — see §20.

Set A — demographic, not used

A1 · Kartik, 26 — tech-savvy, early adopter, automation-friendly. **A2 · Siddharth**, 35 — practical, security-minded, family-first. Neither name appears in any prototype.

A THIRD NAMING SET EXISTS

The prototypes also introduce **Kaustubh** (homeowner), **Arjun** (child) and **Maa** (elder, in agent dialogue). Consistent with Set B's household framing, but the names differ again. If Set B is ratified its names must be reconciled with these, not assumed identical.

THE BINDING CONSTRAINT

P2 is the accessibility floor for the whole design. A low-fluency resident who cannot confirm the air is clean is a product failure. This yields, non-negotiably: body text ≥16pt, touch targets ≥44pt, and no safety-critical state conveyed by colour alone.

6 Jobs to be done

The research names three return-drivers, in this priority order. Everything in v1 serves one of them; a feature that serves none is out of scope.

#	DRIVER	WHAT THE MATERIAL SAYS
1	Checking live AQI	"The #1 stated reason users reopen an air-purifier app — but only if the first number lands as a gut-punch, not a spec."
2	Controlling it remotely	Gives a reason to open the app before you are home. "Start it from the office."
3	Trusting the automations	"Better automations" is the strongest stated driver after AQI itself — and carries the away-mode trap.

7 Product principles

From the name: *ma* (間), the intentional space between things, and *no* (野), nature and openness. The empty space is doing work — it is composed, not left over.

PRINCIPLE	TEST
Whitespace is a feature	If a screen looks sparse that is likely correct. Density is the failure mode here, not a virtue.
One idea per screen	The primary action is findable in one glance.
Calm over informative	Show the state of things, not every metric that produced it. Detail on demand, never volunteered.
Never volunteer alarm	Urgency is earned by real events, not by styling.
Relief, not delight	The emotional target is a glance that confirms everything is fine. Personality lives in small moments, never in the chrome.

Deliberately not

Dashboard-dense · blue-lit · alarm-forward · gamified. And never: make the accent green the ground; require app fluency for basic safety; use a raw metric as a headline; hide a trade-off.

8 The agent model

This is not a feature bolted onto the app — it determines what the app *is*. Home is not a grid of devices; it is a list of what the agent has decided, is proposing, or has already done.

8.1 Three zones, in this order

#	ZONE	CONTAINS
1	Settled now	What is already fine. Per-room readings.
2	Needs your call	Proposals awaiting a decision, badged with a count.
3	Done today	Completed actions, each still undoable.

The ordering is load-bearing: **reassurance, then the ask, then the receipt**. Someone opening the app anxious gets "air is clean in 2 rooms" before they are handed a decision.

8.2 Graduated autonomy — per capability, three levels

The single most important mechanic. Autonomy is **not global and not binary**. It is granted one capability at a time and starts at the bottom.

LEVEL	BEHAVIOUR
Ask me DEFAULT	Agent proposes. Nothing runs without a tap.
Act & tell me	Agent acts, then reports. Undo remains available.
Act silently	Agent acts without notifying.

Capabilities, each with an independent level: `Adjust fan speed` · `Pre-clean before forecasts` · `Reorder filters`.

Right now, I only ever propose. Nothing runs without your tap. As you approve things, you can let me act on my own — one capability at a time.

the agent, in the prototype

TRUST IS PROMOTED ON EVIDENCE

"You've accepted 7 fan suggestions in a row — ready to let this run itself?"

A consecutive-accept count, offered back as the reason. This makes a **per-capability accept/reject tally application state**, not analytics — a real data requirement falling out of a trust decision.

8.3 The explanation contract — four fixed slots

Every proposal exposes "Why this?" with the same four slots, always in the same order. This is a contract, not a copy template.

SLOT	JOB	EXAMPLE FROM THE PROTOTYPE
Sensed	The evidence	"Outdoor AQI at 118 and climbing, forecast 182 by 10pm; bedroom holding at AQI 38 (PM2.5 9 µg/m³)"
Decided	The action	"Pre-clean 8:30–9:00 at medium, then hold windows-closed mode"
Instead of	The counterfactual	"Reacting at 11pm — which would run loud while you sleep"
Obeys rule	The constraint respected	"Quiet hours 11pm–7am stays intact"

WHY THIS MATTERS

"Instead of" is the differentiating slot. Explaining what the agent chose *not* to do is what makes a decision read as considered rather than arbitrary — most assistant UIs ship the first two slots and stop.

"Obeys rule" is a promise: naming the constraint it respected is how the agent proves it is bounded. Ship the disclosure without that slot and standing rules become unverifiable — the user has no way to know their rule held.

8.4 Rehearsal before execution

A proposal previews as a step-by-step timeline, each step individually editable. **Nothing happens until you tap Run.**

TIME	STEP	DETAIL	ACTION
8:30pm	start pre-clean	Bedroom only · medium speed · ~30 min	Edit
9:00pm	windows-closed mode	Recirculate · hold until AQI drops	Edit
11:00pm	respect quiet hours	Fan capped at 28 dB · your standing rule	Locked

Rule-derived steps render Locked, not Edit. A user cannot casually edit away their own constraint from inside a one-off plan. Rules change where rules live.

8.5 Standing rules — goals, not one-time actions

RULE	VALUE	COPY SHOWN
Quiet hours	11pm–7am	referenced by name in "Obeys rule"
Noise ceiling while asleep	28 dB	"Lower is quieter but cleans slower. 28 dB is about a whisper."
Target overnight PM2.5	under 12	"The goal I'll hold the room to, not a one-time action."

Note the copy discipline: each rule states its **trade-off** and anchors its unit to something human. That is part of the model, not decoration.

8.6 Promotion, reversal, isolation

- "Do this always"** promotes an accepted one-off into a standing automation — and appears on *completed* items too, because the moment someone appreciates what it did is the moment to offer the upgrade. The Automation tab then separates *Created by your agent* from *Created by you*.

- **Undo** stays on completed actions, including autonomous ones.
- Recurrence is offered contextually: *"Smog nights repeat this season — do this automatically."*
- The **chat thread is deliberately context-free**: *"Starting fresh — this thread does not carry over anything from Home."* Ambient agent and conversational agent are separate sessions. Anyone reversing this should say why.

THE MOST AMBITIOUS BEHAVIOUR IN ANY SOURCE

The launch prototype shows the agent reasoning **across devices** in one turn — rescheduling the vacuum to 5:30am because yesterday's camera footage showed construction workers arriving at 8am, while answering a question about when Maa usually wakes. That is genuinely impressive and is **the hardest thing to deliver under an on-device constraint** (§16). Nothing in the material addresses how.

Unanswered questions this model raises

- Where does the accept/reject tally live, and does one rejection reset the promotion counter?
- What happens when a proposal's window passes un-actioned? "Later" and "Not tonight" exist; **expiry behaviour does not**.
- Can two proposals conflict, and who arbitrates?
- Is there a declined-consent state?

9 The Week One arc

The most specified thing in the material, and the closest it comes to a product spec. Each day names an emotional target, the research finding, and the screen that delivers it.

D	BEAT	TARGET EMOTION	THE STUDY SAID → SO WE
1	The Reveal	Quiet shock <i>"I didn't know it was this bad."</i>	AQI data is the #1 return reason — but only if the first number lands as a gut-punch, not a spec. → Skip the feature tour. First thing shown is <i>your</i> bedroom air vs the street, colour-coded, already cleaning. First reading: 168 PM2.5 indoors, 211 outside . Day-0 value is a number about you, not an instruction.
2	The Proof	Reassured <i>protection becomes real</i>	Silent protection is indistinguishable from a switched-off box; if we don't narrate the overnight win the value evaporates. → Morning "Air Defended" card quantifies the night: Delhi peaked 187, bedroom never past 22 · 6.5 hours clean · 8.6M particles removed . The ring ties it to the body — SpO ₂ held 97%+ B .
3	The Habit	Hands-off <i>"It just handles it now."</i>	Automations and remote control are drivers #2 and #3. Risk: silent automations make the app feel unnecessary — the away-mode trap. → Three routines built from observed use, one-tap accept: Morning Boost (6:30am, Turbo 20 min) · Sleep Mode (10:30pm, Silent) · Away Saver (on exit, Eco). Plus remote start and a geofence that boosts back 20 min before return.
4	The Neighbourhood	A little proud <i>worth talking about</i>	Comparison and shareability convert a private utility into a daily check-in and an organic acquisition channel — the Spotify-Wrapped lesson. → "Cleaner than 82% of your street" (you 18 · Mehta 2F 54 · Indiranagar avg 96 · outdoor 211) and a one-tap share card. C-10
5	The Worth	Anxiety resolved <i>"It's not eating my bill."</i>	"Is it spiking my electricity bill?" is a quiet usage-suppressor in Indian homes. → Energy card reframes Auto efficiency as money saved: ₹164 saved vs always-Turbo · ₹0.40/hour, less than your ceiling fan — then translates it to 3 trees' worth of clean air . Cost worry becomes a small point of pride.
6	The Gap	Curious <i>"Wait, the rest of the house?"</i>	The first upsell must feel like help, not a pop-up. Filter health is both a retention signal and the consumable lever — surfaced late it becomes a 1-star "it stopped working" review. → Home air map shows what each room still breathes (bedroom 20 · living 94 · kids 61 · kitchen 112 ; "a Living Room unit would cut it to ~24"). Filter health at 34%, ~22 days left , with reorder timed to arrive first. C-9
7	The Score	Sealed in <i>"This is my home's air now."</i>	Day-7 → Week-2 is the canonical retention chokepoint , and the weekly recap is the cheapest acquisition channel a product ever gets. → Air Score 84/100 : 148 hours of clean air · 61M particles removed · cleaner than 82% of Indiranagar · best night Thursday. Shareable, with a Week-2 cliffhanger (pollen / Diwali). C-5

The night-hook → morning-proof pair (D1→D2) is the most important structure in the product.

Night: *"Tonight the AQI outside hits 200+. I'm on Auto in your bedroom — wake up and see what I held off."*

Morning: *"While you slept, Delhi hit 187. Your bedroom never went past 22."* It is what converts a silent box into something worth checking. **Never ship one without the other.**

THE ARC'S SINGLE POINT OF FAILURE

The entire seven days are push-dependent. A notification-permission denial at J-ONBOARD-06 silently kills days 1–7 with no other signal changing. This is the highest-leverage thing to measure (§22) — ahead of the retention number it protects — and it needs a designed no-push fallback, which does not exist.

10 Journeys

Every step below corresponds to a screen that exists in a prototype. IDs are stable; retire rather than renumber. No step carries a persona ID until C-1 closes.

J-ONBOARD · Account and first device

ID	STEP	NOTES
J-ONBOARD-01	Welcome	"By continuing you agree to our Terms & Privacy Policy"
J-ONBOARD-02	Value slides (4)	One App Whole Home · Privacy First (on-device AI, DPDP compliant) · Smart Automations · Home Score
J-ONBOARD-03	Phone number	Country selector, India default
J-ONBOARD-04	OTP verify	Failure here is a hard acquisition leak
J-ONBOARD-05	Name the home	
J-ONBOARD-06	Notification priming	Pre-permission screen shown before the OS prompt. The most consequential step in the product.
J-ONBOARD-07	Pick first device type	Camera · Lock · Purifier · Vacuum only

GAP: slide 4 promises Home Score before any device exists. No zero-device Score state is designed.

J-PAIR-PUR · Add a purifier

The most carefully-designed flow in the material. Three steps, each removing a known failure — and **every step states its purpose, not just its instruction**. Worth preserving as a rule.

ID	STEP	WHAT IT DOES
J-PAIR-PUR-01	Bluetooth discovery	"Power on the purifier. Hold your phone close — this takes a moment." Sets an expectation of delay instead of appearing stuck.
J-PAIR-PUR-02	Wi-Fi, 2.4 GHz	Hides the 5 GHz SSIDs the device cannot join rather than listing them and failing — removing the single most common smart-home setup failure. Says why: "Live air data and automations need a connection."
J-PAIR-PUR-03	Name + room	India-context presets including Pooja Room . Names the reason: "Naming the room lets us compare against the right outdoor sensor."
J-PAIR-PUR-04	First reading	The gut-punch. 168 PM2.5, indoors.

J-AGENT · The proposal loop

ID	STEP
J-AGENT-01	Proposal surfaces in "Needs your call"
J-AGENT-02	User expands "Why this?" — Sensed / Decided / Instead of / Obeys rule
J-AGENT-03	User rehearses the plan as an editable timeline

ID	STEP
J-AGENT-04	Run · Later · Not tonight — expiry undefined
J-AGENT-05	Agent reports back; item moves to "Done today" with Undo
J-AGENT-06	User promotes via "Do this always" → appears in Automation tab
J-AGENT-07	After repeated accepts, agent offers a capability autonomy upgrade

J-REORDER · Consumable reorder

ID	STEP
J-REORDER-01	Filter health surfaces proactively, not at failure
J-REORDER-02	Reorder CTA framed by arrival: "Arrives before it runs out"
J-REORDER-03	Payment · address · confirmation · tracking — DOES NOT EXIST

J-SHARE · Invite a household member B

Pick role → select devices → select per-device features → set window/expiry → send invite → review in audit.

Blocked by C-3 — the flow was rebuilt three times in the prototype (SharingFlow , V2 , V3), which is itself the evidence it is unsettled.

GAP — flows referenced but never designed

MISSING FLOW	CONSEQUENCE
Payment / order	F-SHOP is the recurring-revenue path and stops at a price tag
Offline / device unreachable	No degraded state, despite a notification reading "camera went offline briefly"
Filter fully exhausted	The predictive path is designed; the failure path is not
Declined consent	Face recognition is declinable; nothing shows what a declined camera does
Zero-device Home Score	Promised in onboarding before any device is paired
Proposal expiry	"Later" and "Not tonight" exist; being ignored does not
Multi-home	Data structure exists with no flow — and P1 is <i>defined</i> by managing a second home
Account deletion / data export	A DPDP obligation with no surface
Empty, loading, error states	Absent throughout all four prototypes

11 Feature scope

Thirteen features are registered in `src/hardware/feature-map.json`. Every one exists as a built prototype screen. **None has a PRD.** "Prototyped" and "specified" are not the same thing.

In — v1

F-ID	FEATURE	NOTES
F-AQI-LIVE	Live air quality, per room	Driver #1. Everything else sits on it. Appears in all four sources.
F-AGENT-PROPOSE	Agent proposals + autonomy	§8. Ship at "Ask me" only until C-8 is ruled.
F-FILTER-HEALTH	AQI-weighted filter health	Differentiator <i>and</i> revenue. C-7 — NO FORMULA
F-AUTOMATIONS	Routines, agent- and user-made	Templates: clean before wake · boost when cooking · pause when out. Triggers: schedule, AQI threshold, geofence, wearable, forecast, device event.
F-WEEK-RECAP	Air Week recap + share	Bridges the D7→W2 chokepoint
F-ENERGY	Running cost as saving	Answers the bill anxiety that suppresses usage. Depends on an unstated electricity tariff, which varies by state and slab.
F-AIR-MAP	Home air map	C-9 — SENSORLESS ROOMS
F-SHOP	Filter reorder	Catalogue and price only — no transaction anywhere

Out — v1

F-ID	FEATURE	WHY DEFERRED
F-FACES	Faces You Trust	B Consent flow designed; declined state not
F-HOME-SCORE	Home Score	B C-5 — FOUR DEFINITIONS
F-SHARING	Household sharing	B C-3 — THREE ROLE MODELS
F-NEIGHBOURHOOD	Neighbourhood comparison	C-10 — PRIVACY
F-WEARABLE-BRIDGE	Wearable bridge	C-8 — HEALTH-ADJACENT

Also out: cameras, locks, vacuums, doorbells, lights, music. And **third-party ecosystems (Alexa / Google / HomeKit / Matter) are absent from all four documents** — for a 2026 connected-home product that is a conspicuous silence. Confirm it is a decision, not an oversight.

12 The differentiator

AQI-weighted filter life, not clock-counted hours. Competitors count runtime, so three months in Delhi's winter smog and three months in Mumbai's monsoon return the same filter score.

Every other purifier tells you your filter is 72% healthy — three months in a row. We don't. Because we actually measure it.

Noise Sanctuary

SAME FILTER, SAME HOURS	REAL AQI	COMPETITOR SAYS	NOMA SAYS
Delhi NCR · winter	287	72%	38%
Mumbai · monsoon	94	72%	64%

Identical hardware, identical runtime, **26 points apart**. That gap is the product. A Delhi filter can exhaust in 4 months what a Mysore filter takes 12 months to use up.

Method	AQI-weighted particulate load, computed hourly
Inputs	Live outdoor AQI · indoor particle sensor · fan runtime
Claimed effect	Above AQI 200, a filter absorbs 3–5× more particulate per hour
Alerting	Predictive — surfaces before failure, reorder timed to arrive before exhaustion
Spec status	None. No formula, no coefficients, no validation data

C-7 — THE MOST SERIOUS GAP IN THIS DOCUMENT

This is a **marketing claim awaiting an engineering model**. It is the one thing that makes the product defensible against Xiaomi at the same price, and it is specified nowhere. Closing it is worth more than any other v1 work.

13 Hardware and sensing

Canonical: `src/hardware/devices.json` . All values marked `prototype-only` — nothing verified against a datasheet or BOM.

Purifiers — the anchor product

HW-ID	MODEL	COVERAGE	CADR	FILTRATION
HW-PUR-200	Air Pro 200	up to 280 ft ²	240 m ³ /hr	HEPA H13
HW-PUR-500	Air Pro 500	up to 550 ft ²	400 m ³ /hr	HEPA H13 + activated carbon
HW-PUR-MAX	Air Pro Max	up to 900 ft ²	600 m ³ /hr	Dual HEPA H13

The rest of the ecosystem B

HW-ID	MODEL	SPEC
HW-CAM-360	Cam 360 2K	2K · night vision · PTZ · two-way audio · privacy mode · person & pet detection · cloud + SD storage
HW-BELL-PRO	Bell Pro	No spec supplied
HW-LOCK-X1	Smart Lock X1	Remote unlock · access log
HW-LOCK-X2	Smart Lock X2	Appears in the seeded list only. Difference from X1 unstated.
HW-VAC-R1	Clean Bot R1	LiDAR · 5000 Pa · sweep & mop · auto-empty dock
HW-VAC-R1P	Clean Bot R1 Pro	LiDAR · 8000 Pa · hot-water mop · self-wash + auto-empty

Sensing — the substance behind the core claim

Indoor: on-device **QSensAI laser particle sensor**.

MEASURE	EXAMPLE	NOTE
PM2.5	167 µg/m ³	Fine particles — the headline number
PM10	210 µg/m ³	Coarse dust
NO ₂	48 ppb	Traffic
O ₃	32 ppb	Ozone
VOC	Low	Reported qualitatively, never as a number
Humidity · Temp	44% · 24°C	
CO ₂	480 ppm	Appears in Score factors only. Which SKU carries a CO₂ sensor is unstated C-9

Outdoor: nearest **CPCB** (Central Pollution Control Board) station, *blended* with the on-device reading.

THE BLEND IS SPECIFIED NOWHERE

Every indoor-vs-outdoor number — the product's core emotional beat, and the basis of the entire Week One arc — depends on an algorithm no document describes.

Connectivity

Pairing	Bluetooth LE — device powered on, phone held close
Network	Wi-Fi 2.4 GHz only
Notable behaviour	Setup hides the 5 GHz SSIDs the device cannot join , rather than listing them and failing. This removes the single most common smart-home setup failure and must be preserved.
Firmware	OTA; version surfaced in device settings (prototype: v2.4.1)
Local controls	LED brightness, child lock — must work independent of cloud

Modes

C-4 · THREE INCOMPATIBLE FAN-MODE VOCABULARIES

SOURCE	MODES
Week One doc	Auto · Turbo · Silent · Eco
Agent doc	Auto · Sleep · Turbo · Manual
Launch app	Quiet · Standard · Turbo · Max

"Turbo" is the only value present in all three. No mode UI, automation action or voice copy naming a mode can be finalised until this is ruled.

Distinct and *consistent* across all sources: whole-home scenes **Home · Away · Night**. Vacuum **B**: clean modes Sweep / Sweep & Mop / Mop; suction Eco / Standard / Max / Turbo Max; targeting Whole Home / Select Room / Custom Area / Point Clean.

Wearables **B**

Not NOMA devices — companion hardware, and **signal sources rather than controllable devices**. The integration ties air quality to a body metric, which is the strongest available answer to "why does clean air matter to me tonight?"

SIGNAL	SOURCE	DEFAULT ACTION
WEARABLE_SLEEP B	low activity + HR in sleep range	Night mode
WEARABLE_WAKE B	movement spike after sleep	Morning mode
WEARABLE_SP02 B	blood oxygen below 94%	Purifier → Turbo C-8
WEARABLE_STEPS B	daily steps ≥ 8,000	user-defined
WEARABLE_GPS_LEAVE B	GPS outside 200 m of home	Away mode

SIGNAL	SOURCE	DEFAULT ACTION
WEARABLE_GPS_ARRIVE	GPS within 150 m of home	Home mode

Devices: **Noise Luna Ring 2** (SpO₂, sleep, HR) and a Noise smartwatch (model unspecified).

Consumables

ITEM	PRICE	NOTE
HEPA H13 filter	₹1,299 or ₹1,499	C-6 The highest-frequency purchase in the product
Mop pads ×4 B	₹499	
Cam wall mount B	₹299	
Main brush · side brush B	—	Wear tracked as % (main 72 · side 85 · filter 45 · mop pad 60)

14 Data model

Inferred from the prototypes — no schema, migration or validation exists. Detail: `memory/architecture/data-models.md`.

ENTITY	FIELDS	NOTE
Home	<code>id, name, rooms[], members[], devices[], score, scene(home away night)</code>	Multi-home is implied but has no flow. Whether members, rules and autonomy grants are per-home or per-account is undefined — and P1 is defined by managing a second household.
Room	<code>id, name, icon, devices[], aqi</code>	<code>aqi</code> may exist with no sensor present <code>C-9</code> . Two preset lists exist <code>C-12</code> : 8 items, or 11 including Pooja Room , Bathroom, Front Door, Garage. The 11-item list is the India-appropriate one.
Device	<code>id, model, type, room, name, online, firmware</code> + per-type telemetry	Device name and room are independent fields. The prototype contains "Living Room Purifier" assigned to room "Bedroom" — proving they are decoupled, and that any UI showing one while implying the other will mislead.
Member	<code>id, name, contact, role, grants[], schedule?, expiresAt?</code>	The most developed model in the set. Guest auto-expiry defaults to 24 h. <code>C-3</code>
Automation	<code>id, name, origin(agent user), trigger, actions[], enabled</code>	<code>origin</code> is surfaced in the UI, so it is a real field — the tab splits agent-created from user-created.
Proposal	<code>id, capability, sensed, decided, insteadOf, obeysRule, plan[], state, undoable</code>	All four explanation slots required. <code>plan[].locked</code> marks a rule-derived step. New in v1: expiresAt — currently undefined behaviour.
StandingRule	<code>id, label, value, unit, tradeoffNote</code>	The trade-off copy is part of the model, not decoration.
AutonomyGrant	<code>capability, level(ask tell silent), consecutiveAccepts</code>	New in v1: safetyClass — its absence is exactly the gap in C-8.
Event <code>B</code>	<code>id, type, time, location, identity?, badge, clipAvailable, unread</code>	Badges: Known · Verified · Review
Face <code>B</code>	<code>id, label, embedding, consentGranted</code>	Embedding is on-device only, never uploaded (DPDP Act 2023). Consent explicit, per-feature, revocable.

15 Permissions and sharing `B`

Grants are per device type, and the granularity is real. Deferred from v1, documented because it is the most developed model in the source material.

TYPE	GRANTS	ITEMS
Camera	7	Live view · Basic functions (PTZ, two-way audio, privacy mode) · Events · Cloud recordings <i>view</i> · Cloud management <i>delete</i> · Local SD recordings · Device management
Lock	2	Unlock · View access log
Purifier	2	Control (power, fan, modes) · Schedules
Vacuum	2	Control (start, dock, spot) · Schedules

THE ASYMMETRY IS DELIBERATE AND CORRECT

Camera splits *viewing* cloud clips from *deleting* them, and lock splits *unlocking* from *seeing who unlocked*. A guest may need entry without seeing the household's movements. Preserve this.

C-3 · THREE ROLE MODELS COEXIST

SOURCE	ROLES
ROLES	family · staff (weekly window) · guest (date range)
ROLES3	family · guest (staff folded in)
SHARE_ROLE_BADGE	owner · family · guest

Do not resolve by picking the newest file. Only the 3-role version has an explicit `staff` role — that is persona P3, and collapsing it into `guest` loses the recurring weekly-window case P3 is defined by.

Retention is a billing entitlement, not a storage default: event history is 30 days on the mid tier, 365 on the top tier.

16 Privacy and compliance

Stated repeatedly and specifically across the prototypes. This is India's Digital Personal Data Protection Act 2023 — a legal obligation, not a marketing line.

What the product claims

- **On-device AI.** "DPDP compliant. Your data never leaves home."
- Face data is stored **"only on your device. It is never uploaded to any server."**
- Face tagging is gated by an explicit consent sheet (Decline / Allow & Continue) **citing the DPDP Act 2023 by name.**
- The camera event view is badged **"No cloud · DPDP compliant."**
- The agent is labelled **"On-device · Private"** in the chat header.

C-11 · THE UNRESOLVED CORE TENSION

Three shipped features cannot work locally:

- **Outdoor AQI blending** — requires a CPCB network call
- **Forecast-driven pre-clean** — requires tomorrow's forecast
- **Neighbourhood comparison** — aggregated peer data, server-side by definition

"Never leaves home" and "uses tomorrow's forecast" cannot both be literally true. The claim most likely means *personal and biometric data stays local; ambient public data is fetched* — defensible, and still a strong position. It must be narrowed to that before it appears in store copy or a privacy policy.

Three consequences the prototypes do not resolve

1. **On-device inference bounds model size and latency**, and therefore how clever the agent can be — including the cross-device reasoning in §8.6.
2. **Consent is per-feature, not per-app.** Declining face recognition must leave a working camera. No declined state is designed anywhere.
3. **Telemetry is exfiltration.** Whether analytics is exempt, separately consented, or aggregated on-device before sending is an architecture decision that must be made *before the first event is named*.

Also required and unaddressed: **account deletion and data export** are DPDP obligations with no surface in any prototype.

17 Competitive position

Methods quoted from each brand's own manuals and FAQs (June 2026). Pricing from Amazon India.

BRAND	PRICE	FILTER SCORE BASED ON	AT AQI 300+
Xiaomi Smart Air 4	₹8,999	Hours, speed + sensor	sensor goes blind
Dyson Pure Cool	₹38,900	Mostly a 12-month clock	degrades

BRAND	PRICE	FILTER SCORE BASED ON	AT AQI 300+
Honeywell Air Touch	₹14,990	Hours + speed only	no sensor
Philips Series 1000	₹9,995	WHO regional average	sensor goes blind
NOMA · NS Air S1	₹8,999	Real particle load, hourly	AQI-verified

Independent lab testing (Smart Air, 2026) found Xiaomi and Philips PM2.5 sensors **lose accuracy at high pollution levels — exactly when Delhi's air is most hazardous**. Honeywell's own manual states its indicator "runs on a timer" and may not reflect actual particle load. Philips bases recommended filter life on WHO urban outdoor averages.

THE PRICE PARITY IS THE WHOLE ARGUMENT — AND IT IS CONTESTED

The deck positions NOMA at **₹8,999, matching Xiaomi's entry price**, and the entire competitive case rests on "same price, honest filter score." The launch app prices the same product at **₹4,999** **C-6**. Nearly 2×. One of the two documents is describing a different product; this is not a typo to smooth over.

Claimed exclusives beyond filter honesty **B** : full home graph (vs "app only") · family-aware (kids, elders, pets, grandparents) · camera + vacuum integration · a single whole-home Score.

18 Monetization

Hardware

SKU	PRICE	TAG
Cam 360 2K B	₹2,999	Bestseller
Bell Pro B	₹3,499	
Air Pro	₹4,999 or ₹8,999 C-6	
Smart Lock X1 B	₹5,499	New
Clean Bot R1 B	₹8,999	Popular

Consumables — the recurring line

HEPA H13 filter at ₹1,299 or ₹1,499 C-6. **Filter health is the strategic centre of the product:** simultaneously the retention signal, the differentiator, and the consumable trigger.

Surfaced too late it becomes a 1-star "it stopped working" review. Surfaced predictively it becomes a reorder that arrives before it runs out.

the same data, opposite outcomes, decided by when it is shown

Subscription B

Deck only — **no paywall, upgrade prompt or gated feature exists in any app prototype.**

TIER	PRICE	INCLUDES
Free	₹0 forever	Live control & room views · real-time AQI in every room · real filter health score · current Score · basic automations
Plus MOST POPULAR	₹699/mo ₹6,999/yr	AQI-weighted filter tracking · predictive replacement alerts · doorstep filter cleaning (3+ purifiers) · 30-day event history · pet detection · family awareness
Family	₹1,199/mo ₹11,999/yr	All Plus · elder monitoring & check-ins · routine intelligence & alerts · 365-day home memory · annual home health report

C-13 · THE TIERING CONTRADICTS THE POSITIONING

Free gets "real filter health score"; Plus sells "AQI-weighted filter tracking". **But AQI-weighting is what makes the score real** — it is the entire differentiator versus the competitors the deck names. Selling it as an upsell undercuts the headline claim, and the two tier descriptions cannot both be honest. (Also: annual pricing implies "save 2 months" on Plus but only ~1.6 on Family, stated as if identical.)

Contextual upsell — the designed surfaces

Monetisation is **data-triggered, never a standalone pop-up**. The one thing all sources agree on. Maps directly to persona P4.

SURFACE	TRIGGER	ASK
Home air map	Rooms without a purifier show what they still breathe (living 94, kitchen 112)	A second device — "a Living Room unit would cut it to ~24"
Filter card	Predicted exhaustion (~22 days)	Reorder
Shop reorder rail	Device already owned	Matched consumable
Week recap	Day 7	Retention, then the Week-2 hook

The stated rule: "the first upsell has to feel like help, not a pop-up."

A TOKEN-LEVEL CONSEQUENCE OF A BUSINESS DECISION

Every upsell surface wants an accent-coloured CTA — and the accent budget is 10% of a screen with **one** primary action (\$20). Upsell and primary action therefore compete for the same budget. On the air map the upsell *is* the primary action; on the filter card it is not. Unresolved anywhere.

Acquisition, treated as a revenue line

- "Every share is a billboard on a friend's phone."
- "The weekly recap is the cheapest acquisition channel a product ever gets."
- **Content rule, learned the hard way:** "Pre-filled brand captions kill shares; emoji + the number travel."

Gated on C-10 — the share engine depends on neighbourhood comparison, which has an open privacy question.

NOTHING EXISTS BEHIND ANY PRICE

No payment gateway, address book, tax handling, invoice, refund, subscription management or order tracking appears in any prototype. **Every number in this section is a price tag with nothing behind it.**

19 Voice and content

The one rule everything follows

Observed → acted → what it means for you. Never a bare status.

Good morning. While you slept, Delhi hit 187. Your bedroom never went past 22.
the morning-proof notification

Three moves: what happened outside, what I did, what you got. A status line would have said "Air quality: Good." **The difference is the entire product.**

Voice characteristics, with evidence

TRAIT	EVIDENCE
First person, owns the action	"I'm on Auto in your bedroom" — not "purifier switched to Turbo"
Pairs a number with a consequence	"₹0.40/hour — less than your ceiling fan"
Anchors units to something human	"28 dB is about a whisper"; filter life as "~22 days left", never a bare %
States the trade-off	"Lower is quieter but cleans slower"
Pre-empts a known anxiety	The energy card exists because bill worry suppresses usage
Asks before autonomy, with evidence	"You've accepted 7 fan suggestions in a row"
Warm without being cute	"Heading out? I'll stop wasting power." No mascot, no exclamation marks, no apologising.

Notification patterns

TYPE	EXAMPLE	JOB
Night hook	"Tonight the AQI outside hits 200+. I'm on Auto in your bedroom — wake up and see what I held off."	Creates a reason to open the app tomorrow
Morning proof	"Good morning. While you slept, Delhi hit 187. Your bedroom never went past 22."	Pays off the hook
Explained action	"Turbo kicked in — outdoor AQI spiked to 167"	Narrates autonomy so it is not invisible
Care	"Filter at 82% · about 48 days of clean air left"	Retention + revenue, framed as care
Social	"You've stayed in the cleanest 20% of Indiranagar all week. Share it?"	Acquisition

The share-caption rule

The one place brand voice must disappear entirely. A share card speaks as the **user**, not the company:

Outdoor was 211 today. I'm breathing 12× cleaner air. 🌿

the share card, in the user's voice

India context is a correctness requirement, not flavour

SPECIFIC	WHY IT IS LOAD-BEARING
Delhi 287 winter vs Mumbai 94 monsoon	The contrast <i>is</i> the filter argument
₹ and ₹/hour vs a ceiling fan	The universally understood cost baseline
Pooja Room as a first-class preset	A room-preset list without it is wrong for this market
"Maa" in agent dialogue	Natural household register
Blinkit as a recognised delivery identity	Camera events read as local life
CPCB as the outdoor authority	Not WHO averages
Diwali and pollen season	Scheduled air events, used as the Week-2 hook
2.4 GHz-only routers	An Indian-router reality the setup flow designs around
Household staff as a normal role	Not an edge case — persona P3

THE ASSISTANT HAS NO SETTLED NAME — C-2

Four names for one agent: **Ding** ("🔔 Ding now" in Week One) · **Noise AI Agent** (launch app) · **Sanctuary Assistant** (deck) · *unnamed first person* (Air Agent). The product is now NOMA, so none is currently correct. **Every string above is unshippable until this is ruled.**

O-5 · A HINDI LOCALE EXISTS AND CANNOT BE TYPESET

`src/i18n/locales/hi/` exists, the primary personas are Indian, and **Google Sans Flex has zero Devanagari coverage** — confirmed by reading its `cmap`. Every string in this section is English-only. A paired Devanagari face is required before `hi` ships.


20 Design system

The one area with no open conflicts. ADR-001...005 are settled. Canonical:

`src/tokens/design.tokens.js`.


Surfaces — warmth lives in the ground, not the greys

Derived from Airbnb iOS (observed at 393×852pt). Two ramps, deliberately separated: `neutral` carries text and structure, `sand` carries the warmth. Tinting the text greys too would drift contrast ratios and muddy photography.



The ramp stops at `#222222` for text. Pure black is reserved for system chrome — it reads harsh against a warm ground. Airbnb's Rausch red was explicitly rejected.

Accent — two greens, budgeted at 10%



Green occupies **~10% of any screen, never more**. It marks the one thing that matters; it is never the ground. Emphasis that is *not* the primary action stays near-black — that is what protects the budget. The pair sits 16° apart in hue (neon 93°, moss 109°), so they read as one green family, light and dark.

Measured contrast — one pairing is a trap

PAIRING	RATIO	VERDICT
<code>accent.solid</code> + white — moss on white	6.27:1	AA correct for a 16pt/600 button label, but not AAA
<code>accent.bright</code> + <code>#222222</code> — neon on ink	8.66:1	AAA correct
neon on white	1.84:1	FAILS never ship

Neon is the trap: it is light enough to look like a coloured button that should carry white text, and it is unreadable. Always pair a green with the `on*` token beside it.

Cards — two anatomies, not alternatives

	CARD ELEVATED	CARD OUTLINED
Sits on	warm ground (sand.200)	white canvas
Radius	20	12
Defined by	soft two-layer shadow	1pt hairline
Use for	hero, featured, identity	dense structured content

Never combine a hairline border with a shadow on the same card. **Never** place an elevated white card directly on white — without the warm ground there is nothing to lift off, and the shadow reads as dirt.

Shape, spacing, elevation

Radius	sm 8 · md 12 — default · lg 16 · xl 20 (elevated cards, sheets) · full 999 (pills)
Spacing	4pt base — 4 · 8 · 12 · 16 · 20 · 24 (screen gutter) · 32 (section gap) · 40 · 48 · 64
Card padding	16
Elevation	Every level is two layers — a tight contact shadow plus a wide ambient one — at 4–6% opacity. subtle · card (default) · floating · onMedia (20%, for controls over photography). A single hard drop shadow is never correct.

Typeface — Google Sans Flex

Variable, SIL Open Font License, owner-supplied. Six axes; NOMA drives wght and opsz and pins width , slnt , GRAD , ROND .

opsz is the axis worth wiring up: optical size should track rendered size — low opsz thickens strokes and opens spacing for small text, high opsz refines them for display. Every step carries its own opsz so this survives into engines without font-optical-sizing: auto .

STEP	SIZE / LEADING	WEIGHT	USE
display	32 / 36	700	Screen title
title1	26 / 31	700	
title2	22 / 27	600	Section, big figures
title3	18 / 23	600	Card title
body	16 / 22	400	P2 accessibility floor
bodyStrong	16 / 22	600	Buttons
caption	14 / 19	400	Supporting detail
micro	12 / 16	500	Badges, timestamps

OPEN — NEEDS A TOKEN DECISION

The scale has no hero-numeral step. The device AQI readout needs ~56pt and the scale tops out at 32. Currently used as a flagged one-off; it wants a ratified typography.scale addition.

Open design items

O-3 light theme only — tokens are semantic so a dark map can be added without touching UI code, but no dark values exist. **O-5** no Devanagari (§19). **O-6 no semantic status colours — green cannot carry "bad"**, which blocks any real device-status UI. **O-4** hexes and radii were read optically from captures (cross-origin sampling was blocked); close but re-verify on device before locking a 1.0.

21 Motion

ADR-005. Canonical: `src/tokens/motion.tokens.js`. Gates: `.claude/rules/motion.md`.

Bounce means "something happened"; stillness means "everything is fine."
the rule that resolves playful-vs-*ma*

Motion is **playful — springy, bouncy, Airbnb-register — on discrete moments** (a tap, an arrival, a confirmation) and **almost perfectly still on ambient ones** (live readings, monitoring, charts). Calm is the ground; bounce is the event. That split is how owner-stated playfulness and the *ma* philosophy stop fighting.

TOKEN	VALUE	USE
<code>durations</code>	80 · 140 · 240 · 360 · 480 ms	hover → tap → card → expander → sheet
<code>durations.ambient</code>	4000 ms	The only long value. Breathing ripples, idle pulses
<code>easings.spring</code>	0.34, 1.56, 0.64, 1	~10% overshoot — the house spring. Buttons, chips, toggles, cards
<code>easings.springBold</code>	0.22, 1.8, 0.36, 1	Hero arrivals — sheets, overlays, toasts
<code>easings.exit</code>	0.4, 0, 1, 1	Every exit. Never bounces
<code>stagger</code>	45 ms lists · 60 ms grids	Between siblings entering
<code>scale</code>	press .96 · pop 1.04 · breathe 1.02	<code>breathe</code> is the ambient ceiling

Registered patterns

`M0-PRESS` press & release · `M0-POP` arrival pop · `M0-STAGGER` list entrance · `M0-SHEET` sheet rise · `M0-KNOB` toggle thumb · `M0-TOAST` · `M0-PULSE` count change · `M0-BREATHE` ambient (the only looping one).

Non-negotiable gates

1. **Tokens only** — UI code never contains a raw duration or bezier.
2. `prefers-reduced-motion` is **mandatory** on every animation. An accessibility gate, no exceptions.
3. **Exits never bounce.** Overshoot on the way out reads as broken, not playful.
4. **Ambient never springs.** A live number does not bounce because it changed.
5. **Alarm never bounces.** If something is wrong, motion gets out of the way.

PROVENANCE NOTE

The owner referenced a bouncy style defined in an earlier session unavailable to this one. These values are a fresh codification of the stated direction (playful / fun / quirky / bouncy / Airbnb-like). If the earlier values surface, reconcile by revising ADR-005 — not by editing UI code.

22 Analytics and success metrics

THERE ARE ZERO ANALYTICS EVENTS

All four prototypes were searched for `track( , logEvent( , analytics , gtag , amplitude , mixpanel` — no matches. **The product has a fully articulated seven-day retention strategy with named chokepoints and stated behavioural mechanisms, and no way to measure any of it.**

`src/analytics/events.schema.json` stays empty deliberately: a guessed machine-readable taxonomy looks authoritative and is worse than an empty file. **Blocked by C-11** — the on-device privacy claim must be reconciled with telemetry before the first event is named.

Priority order once unblocked

#	METRIC	WHY IT RANKS HERE
1	Notification permission granted / denied at <code>J-ONBOARD-06</code>	The whole Week One arc is push-dependent. A denial kills D1–D7 invisibly — no other signal changes. Ranks above the retention number it protects.
2	D7 recap seen → returned in week 2	<i>The</i> retention number. Every other Week One measurement is diagnostic for this one.
3	Agent accept / defer / decline / undo per proposal	Application state that doubles as the best available measure of whether the agent is trusted
4	Undo rate per autonomy level	The honest signal that "Act silently" was granted too early. Nobody has asked for this yet.
5	"Why this?" expansion rate	Whether the explanation contract is load-bearing or ignored
6	Filter surfaced → reorder tapped → purchased	The revenue funnel. Cannot close today — no payment flow exists.
7	2.4 GHz connect success at <code>J-PAIR-PUR-02</code>	Validates the 5 GHz-hiding decision, which currently rests on an unmeasured claim

Hard constraints on any taxonomy

1. **No face, identity or biometric value may ever be an event property** — not a name, not an embedding, not a per-individual known/unknown flag.
2. Neighbourhood comparison must not be instrumented until C-10 resolves.
3. Wearable SpO₂ and sleep are **health data**. Measuring *that an automation fired* is not the same as transmitting the reading — keep that line explicit.
4. Every registered event must cite a `J-*` step.

23 Risks

SEV	RISK	MITIGATION
CRITICAL	C-7 — the differentiator has no engineering model	Get a filter-decay spec before building F-FILTER-HEALTH . Do not ship the claim without it.
CRITICAL	C-11 — privacy claim vs three networked features	Narrow the claim to personal/biometric data before any store copy or policy. Unblocks analytics.
CRITICAL	C-8 — no safety class for autonomy	Ship "Ask me" only. Classify capabilities before enabling silent action or unattended purchase.
CRITICAL	C-2 — five product/assistant names	Rule the name. All copy, i18n, prompts and SKUs are downstream.
HIGH	Push denial silently kills the retention arc	Measure it first (§22 #1); design a no-push fallback for D1-D7
HIGH	No payment flow	Recurring revenue is unreachable until it exists
HIGH	Zero error / offline / empty states anywhere	Specify per feature; the PRD template now requires them
HIGH	O-6 — green cannot signal "bad"	Add semantic status colours before any device-status UI
HIGH	C-6 — the anchor product has two prices	Resolve before any pricing surface; the competitive argument depends on ₹8,999 parity
MEDIUM	O-5 — no Devanagari for the hi locale	Pair a Devanagari face before hi ships
MEDIUM	Prototype specs unverified	Confirm CADR, coverage and pricing against a datasheet / BOM
MEDIUM	Sources live in ~/Downloads	Vendor the prototypes into docs/prototypes/
MEDIUM	Cross-device agent reasoning may be undeliverable on-device	Scope the agent's reach explicitly in the F-AGENT-PROP0SE PRD

24 Open conflicts

Nothing here is resolved. Where sources disagree, both readings are recorded with evidence. Full detail: `memory/decisions.md`.

ID	SEV	CONFLICT	BLOCKS
C-1	HIGH	Two persona sets. Prototypes corroborate Set B four ways.	personas.json · tone routing · persona IDs on journeys

ID	SEV	CONFLICT	BLOCKS
C-2	CRIT	Product + assistant have five names	all copy · i18n · prompts · glossary · SKUs
C-3	HIGH	Three role models; flow rebuilt three times	F-SHARING
C-4	MED	Three fan-mode vocabularies; only "Turbo" is shared	mode UI · automation actions · mode copy
C-5	HIGH	Four Score definitions — only one weighted (Air 25 / Security 25 / Cleanliness 20 / Filter 15 / Energy 15), only one with a Family dimension, one possibly a different feature	F-HOME-SCORE
C-6	HIGH	Purifier ₹4,999 vs ₹8,999; filter ₹1,299 vs ₹1,499	all pricing · the competitive argument
C-7	CRIT	Differentiator has no formula	F-FILTER-HEALTH — differentiator <i>and</i> revenue
C-8	CRIT	No safety class for autonomy; SpO ₂ trigger is health-adjacent	all autonomy above "Ask me"
C-9	HIGH	Air map shows AQI for sensorless rooms — and sells hardware on it	F-AIR-MAP credibility
C-10	HIGH	"Anonymised" comparison shows a surname + floor ("Mehta 2F")	F-NEIGHBOURHOOD · sharing · acquisition
C-11	CRIT	"Data never leaves home" vs three networked features	analytics taxonomy · store copy · privacy policy
C-12	MED	Two room-preset lists; only one has Pooja Room	onboarding · room creation
C-13	HIGH	Subscription sells the free tier's own claim	pricing page · tier copy

25 Sequencing

Before any build — decisions only

1. **The §3 scope question** — which product ships first
2. **C-2** — the name
3. **C-7** — the filter-decay model
4. **C-11** — privacy vs telemetry
5. **C-8** — autonomy safety classes

Then, in order

F-AQI-LIVE → F-AGENT-PROPOSE at "Ask me" → F-FILTER-HEALTH → F-AUTOMATIONS → the D1/D2 push pair → F-WEEK-RECAP → F-SHOP *with a real payment flow*.

In parallel

- Per-feature PRDs from `templates/feature-prd.md` — each requires error, empty and offline states
- The analytics taxonomy, once C-11 clears
- **O-6** semantic status colours — nothing can report "bad" until these exist
- A Devanagari pairing for `hi`
- Vendor the source prototypes into the repo
- Verify hardware specs against a datasheet

26 Declares

Mandatory block — this is what makes the PRD machine-checkable against the canonical registries.

Data

`Home` · `Room` · `Device` · `Automation` · `Proposal` · `StandingRule` · `AutonomyGrant` .

New fields required by v1: `AutonomyGrant.safetyClass` (C-8) and `Proposal.expiresAt` (currently undefined behaviour). Detail: `memory/architecture/data-models.md` .

Events

None defined. Blocked by C-11. Seven required metrics are specified in §22 and deliberately left unnamed so nobody ships them as a taxonomy. Canonical home when unblocked: `src/analytics/events.schema.json` , each event citing a `J-*` step.

Hardware

v1: `HW-PUR-200` · `HW-PUR-500` · `HW-PUR-MAX` · `HW-CONS-HEPA` .

Capabilities required: `indoor-particle-sensor` · `outdoor-aqi-feed` · `forecast-feed` · `standing-rules-engine` · `audit-log` · `scheduler` · `payments` (*does not exist*).

Deferred to B : `HW-CAM-360` · `HW-BELL-PRO` · `HW-LOCK-X1/X2` · `HW-VAC-R1/R1P` · `HW-WEAR-RING2` · `HW-WEAR-WATCH` .