

The NOMA Voice

How the app talks, everywhere — and why. A voice for a home that's looking after you, rather than a wizard processing you.

Status: draft — the register is settled app-wide, the speaker is not.

Every "I" in this document rides on C-2 , still open: whether the product and the assistant are the same character. Treat every example as the register to hit, not the final shipped string.

VERSION

0.1 · 2026-08-07

CANONICAL FOR

tone-matrix.md · glossary.md

BLOCKED BY

C-2 — the name

SCOPE

App-wide

EXTENDS

ai-personas.md

WORKED EVIDENCE

first-run/prd.md §8

§ Contents

1. Why this exists	8. What stays literal, on purpose
2. What makes a setup flow feel like a conversation	9. Errors, specifically
3. The seven rules	10. Before / after — the actual rewrite
4. Who is speaking — the open question	11. Localization
5. Grammar and mechanics	12. The checklist
6. Vocabulary — words in, words out	13. Where this applies beyond onboarding
7. Tone by moment	14. Cross-references

1 Why this exists

The first build of onboarding copy was correct and lifeless. Every screen followed the same shape — a headline stating what to do, a line of body text stating why, a button labelled with the action. Nothing was wrong with it. Nothing in it sounded like anyone.

The owner's instruction was specific: **stop writing instructions, start writing a conversation**. The app should talk about the house, refer to itself as "I", and sound like it's genuinely interested in getting this right for you — not administering a checklist.

This document is the rule set that came out of that instruction. It is not a copy deck — the actual strings live in `docs/features/first-run/prd.md` §8 and in the prototype's source. This is the *why* behind them, written so the next hundred strings can be judged against the same standard without re-deriving it.

2 What makes a setup flow feel like a conversation

A wizard and a conversation can carry identical information and still feel like opposite relationships. The difference sits in two structural habits, not in word choice.

A question and its reason arrive in the same breath

If a screen asks for something, the next sentence — not the help text, not a linked policy — says why. *"Where does it live?"* followed immediately by *"So I can compare your air against the right patch of outdoors."* Left to be inferred, or deferred to a privacy policy, the same ask reads as a form.

Nothing is asked before the app has shown it understands the situation

An ask that lands cold, with no context, reads as a form filling itself out at you. An ask that follows an observation about what's actually in front of the user — a device that just connected, a room that needs naming — reads as something a person would say next, not a step in a sequence.

Neither habit is decorative. Both are trust-sequencing: the app earns the right to ask by showing it already understands, before it asks.

3 The seven rules

Every rule below has a NOMA line already shipped in the prototype — these aren't aspirational, they're a description of work already done, so they can be checked against the artifact.

1 First person, singular

The app is "I", not "we", not "the app". *"We'll text you a code"* becomes **"I'll need Bluetooth for a minute."** A team behind a curtain is a company; a single voice is something you can have a relationship with.

2 Ask, then say why, in the same breath

The reason is never optional, never deferred, never longer than the ask. **"Where does it live?" → "So I can compare your air against the right patch of outdoors."** If the reason isn't in the next sentence, the screen fails this rule regardless of its adjectives.

3 Take the anxiety out of the answer

Setup questions carry tiny stakes that don't need to be there. **"Only really matters once there are two of them"** about naming a device tells the user the question in front of them carries almost no weight, so they can answer without overthinking.

4 "Let's" for shared work

Anything framed as something both parties are doing reads as collaborative. **"Let's get the filter out of its bag"** is a shared task; "remove the filter from its packaging" is a work order. Same information, opposite relationship.

5 Privacy as a concrete negative, not a policy

Never "your data is protected" — always the specific thing that does not happen, stated plainly enough to picture. **"I only notice you crossing in and out. I never keep a trail of where you've been."** A concrete negative is falsifiable, and therefore trustworthy.

6 Warm buttons at warm moments

"Yes, do that" reads as an answer in a conversation. "Confirm" reads as UI chrome. Reserve the warmest buttons for moments the user is glad to be there; keep physical/technical steps a notch more plain, because those need to be found fast, not savored.

7 Observation, then offer, then ask

Never ask before demonstrating understanding. On Home: **"It's 168 out there. I've already started."** — the observation and the action, before anything is asked at all.

4 Who is speaking — the open question

Every "I" in this app is written as if the product itself is speaking. That is a **product decision wearing a copy decision's clothes**, and it rides on **C-2**, which is still open.

`memory/voice/ai-personas.md` already established this exact register for the ongoing product — first person, owns the action, never a bare status — but under a name that was never settled (four different names for the assistant appear across the source material). This document runs the same voice everywhere in the app, on the assumption that the thing greeting you at the door and the thing telling you the filter's dying are the same speaker.

IF THEY ARE NOT THE SAME SPEAKER

If the product is NOMA and the agent is a separately named character that lives inside it, every "I" in this flow needs re-attribution before it ships. That is not a global find-and-replace: it changes whether the door, the setup wizard, and the agent's Day-2 proposals are one relationship or three introductions.

Resolve **C-2 before any string in this document becomes an i18n key.** Until then, treat every example here as the register, not the final line.

5 Grammar and mechanics

Rules a writer or a reviewer can apply mechanically, independent of judgment calls about warmth.

- **Contractions, always.** "I'll", "it's", "you're" — never "I will", "it is". A voice that doesn't contract reads as formal no matter what it's saying.
- **First person for the app, second person for the user.** "I'll need Bluetooth", "your purifier". Never third person for either.
- **Present tense for state, past tense for action taken.** "It's on Sharma_Home" (state); "I've already started" (action). Don't blur the two.
- **No exclamation marks.** Calm is the register; enthusiasm punctuation undercuts it. Warmth comes from word choice, not punctuation.
- **Sentence-case headlines, never Title Case.** "Which one did you bring home?" not "Which One Did You Bring Home?" — Title Case reads as a form label.
- **Numerals, not spelled-out numbers** — except a felt duration read aloud ("about twenty minutes") beats a measurement ("about 20 minutes"); a displayed number stays numeral.
- **Em dashes for the aside, not the semicolon.** Conversational rhythm uses dashes; semicolons read as written, not said.
- **One idea per sentence.** If a sentence needs a comma to hold two instructions, it's two sentences. Read every line aloud — if you'd need to breathe mid-sentence, split it.

6 Vocabulary — words in, words out

Canonical source for `memory/voice/glossary.md` 's banned/approved list.

BANNED (WIZARD VOICE)	USE INSTEAD	WHY
“Please enter your...”	“What’s your...” — a direct question	“Please enter” is a form filling itself out at you
“Invalid [X]”	Name what’s actually wrong, plainly	“Invalid” blames the input, not the mismatch
“Error” / “Failed”	Say what happened and what still works	Both are a verdict with no next step
“Please wait...”	Name the wait and its length	An un-named wait reads as a stall
“Continue” on a warm moment	A warmer verb — reserve “Continue” for plain steps	The mistake is warming the wrong buttons, not using “Continue” at all
“Kindly” / “Please note that”	Cut it — say the thing	Both are throat-clearing with no content
“In order to”	“To”	Three words doing one word’s job
“Utilize”	“Use”	—
“We are unable to”	“I can’t”	First person, direct, no corporate distance
“The user” / “the device”	“You” / “it” (named if ambiguous)	Third person here is the product discussing the person as if they’d left the room
“Successfully completed”	Say what’s true now	“Successfully completed” is a system log line, not something a person says

TWO WORDS WORTH NAMING

"Let's" (Rule 4) and **"I'll"** (Rule 1) carry real weight. If a screen has neither a "let's" moment nor an "I'll" statement, check whether it needed a person speaking at all, or whether it should stay plain (§8).

7 Tone by moment

Not every screen wants the same amount of warmth. Each row is the register for one moment, with a real line already in the prototype.

MOMENT	TARGET	DO	DON'T
Arrival	Anticipation, not friction	"Let's make your home breathe easier"	A brand slogan detached from the moment
Asking for something	Low-stakes, reason up front	"Want the air sorted before you get home?"	A form label with a tooltip
Waiting	Named, bounded, released	"Fifteen seconds or so"	A bare spinner, or "Please wait"
Something went wrong	Device is the subject; offer the next move	"I can't see it yet. Usually it just needs a small nudge."	"Invalid", "Error", or blaming the user
Success	Warm, specific, closes the loop	"That's the hard part done"	"Setup complete." with nothing else
Agent narrating itself	First person, evidence over claims	"It's 168 out there. I've already started."	Marketing copy about the product in third person
System dialogs	Not ours to write	Leave the OS's register alone	A chatty system dialog — reads as a fake one

8 What stays literal, on purpose

Two categories are deliberately not warmed up — a judgment call worth defending, not an oversight.

The instruction inside an instruction screen

The frame around a physical step can be warm; the step itself has to stay literal, because getting it slightly wrong costs the user real time with a purifier in pieces on the floor. "Let's get the filter out of its bag" is warm framing; underneath it is still a precise statement of a failure mode. Unwrapping the filter is the single most common invisible setup failure in the whole flow — charm is not worth the ambiguity there.

System dialogs

iOS and Android own that register. A permission sheet written in NOMA's voice reads as a spoofed system dialog, which is worse than a cold one — it looks like the app is trying to sound official rather than being official.

9 Errors, specifically

Errors are the hardest case: the app has to stay warm and stay unambiguous, and those two goals fight harder here than anywhere else in the flow. Four rules, non-negotiable.

1. **The device is the subject, never the user.** "I can't see it yet", not "you didn't hold your phone close enough." The user did not cause the failure; do not imply they did.
2. **Every error offers the next move, not just a verdict.** "Worth another look?" — never a bare restatement of what failed.
3. **Never apologize on behalf of hardware.** "Sorry, something went wrong" is empty. State what happened and what still works instead.
4. **The fix stays literal even when the frame is warm.** "That password didn't work. Worth another go?" is warm framing around a precise, unambiguous fact.

10 Before / after — the actual rewrite

Ten screens from the shipped prototype, as they were and as they are now. This is the concrete evidence the rules above describe.

SCREEN	BEFORE	AFTER
Get started	“Quiet, clean air. One app.” / [Get started]	“Let’s make your home breathe easier.” / slide to enter
Phone number	“What’s your number?”	“First — how do I reach you?”
Profile	“Who are we setting this up for?”	“And who am I looking after?”
Email verify	“Check your email”	“Have a look in your inbox”
Choose purifier	“Which one is it?”	“Which one did you bring home?”
Plug in	“Plug it in and switch it on”	“Now find it a socket”
Bluetooth	“Switch on Bluetooth”	“I’ll need Bluetooth for a minute”
Connected	“Your purifier is set up”	“That’s the hard part done”
Which room	“Which room is it in?”	“Where does it live?”
Not found	“We couldn’t find your purifier”	“I can’t see it yet”

11 Localization

This voice raises the translation bar. Idiomatic phrases — “a spare key”, “hum away happily”, “already spoken for” — carry meaning through cultural reference, not through their literal words. A machine or literal translation of any of them into Hindi will read as broken, not warm.

THIS NEEDS TRANSCREATION, NOT TRANSLATION

Budget a native-voice pass that re-derives the *feeling* of each line in Hindi, rather than translating the English sentence. The seven rules in §3 should hold in Hindi even where the specific idioms don't survive.

hi is required for this flow, not deferred — memory/voice/ai-personas.md already flags that a resident persona with low app fluency is exactly the case an English-only setup fails. This voice makes that requirement more expensive, not less necessary.

12 The checklist

Before a new string ships, run it against these eight questions.

1. Would a person actually say this out loud to someone in their home?
2. Is the reason for the ask in the same sentence or the very next one?
3. If this is an error, is the device the subject — not the user?
4. Does every error offer a next move, not just a verdict?
5. Is there a wait here that isn't named and bounded?
6. Did this sentence need a comma to hold two instructions? If so, split it.
7. Is this a physical/technical instruction dressed in so much warmth that it's gotten ambiguous? (§8 — the frame can be warm; the instruction can't blur.)
8. Is this a system dialog? If yes, leave it in the OS's own voice.

A string that fails any of these isn't necessarily wrong, but it needs a reason to survive the exception, and that reason should be written down next to it — as every deliberate exception in §8 is here.

13 Where this applies beyond onboarding

This document generalized from an onboarding-only PRD to an app-wide one on 2026-08-07, because the register it describes was never actually specific to setup — a wizard-vs-conversation problem shows up anywhere the app has to ask something, wait for something, or report back. The seven rules, the grammar in §5, and the banned words in §6 apply without modification to every surface below. What changes by surface is *how much warmth a moment can carry* (§7 already covers this axis) — not the rules themselves.

Reconciling with `ai-personas.md` . That document already specified the ongoing-product voice — first person, owns the action, pairs a number with a consequence, anchors units to something human, never volunteers alarm — derived independently from PM prototype material for notifications and Home. It does not conflict with anything here; it's a *superset* for those two surfaces, with rules specific to what an ambient, always-on voice needs that a one-time setup flow doesn't (the night-hook/morning-proof pairing, the share-caption exception where brand voice must disappear entirely). Read both for Home and notifications; this document alone is sufficient everywhere else.

Surface by surface — what's rewritten, and what's still owed

SURFACE	STATUS	NOTES
Onboarding & device setup	Rewritten and shipped.	The worked evidence throughout this document. <code>first-run/prd.md</code> §8.
Home & notifications	Already this register, via <code>ai-personas.md</code>	Written independently but consistent with the rules here. Extend, don't duplicate.
My Home (six-state dashboard prototype)	Not yet rewritten.	<code>docs/features/my-home/build-home.py</code> has its own copy, written before this voice existed. Next concrete pass.
Sharing / People, Automations, Shop, Settings	No prototype exists yet.	Apply the checklist (§12) at first-draft time rather than retrofitting later — cheaper than a second pass.
Errors app-wide	§9's four rules apply everywhere, not just onboarding's error states.	The device (or the automation, or the routine) stays the subject; every error still offers a next move.

DO NOT MARK A SURFACE DONE FROM BELIEF

A surface isn't in this voice because someone believes it is; it's in this voice because a specific line can be pointed at, the same way §10's before/after table checks onboarding.

14 Cross-references

<code>ai-personas.md</code>	The ongoing-product voice for Home and notifications, reconciled with this document in §13. Read it first if <code>C-2</code> closes and every “I” needs re-attribution.
<code>tone-matrix.md</code>	This document is canonical for all its rows , not just onboarding. The stub should point here rather than duplicate this content.

`glossary.md`

This document is canonical for §6’s banned/approved words.

`first-run/prd.md` §8

The onboarding copy this document explains, screen by screen — the only surface fully checked against §10’s before/after standard so far.

`decisions.md` C-2

The open blocker on who “I” is.

The NOMA Voice · v0.1 · 2026-08-07 · generalized app-wide the same day it was compiled from the first-run prototype's shipped copy. Every rule here is checked against a real, already-written line, and §13 tracks exactly how far the rewrite has actually reached. Canonical markdown: `docs/voice/voice-prd.md`.